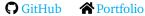# Aniket Saxena

✉ aniketgenai1@gmail.com   GitHub   ⌂ Portfolio

With over 6 years of experience in:

- Proven expertise in Python, machine learning, deep learning, Natural Language Processing (NLP), Computer Vision (CV), and advanced generative AI techniques, with a strong focus on developing innovative and scalable solutions.
- Leadership in driving cross-functional teams of data scientists and engineers, delivering production-ready machine learning systems, collaborating with stakeholders and clients to define project scope, success criteria, and deployment strategies.
- Dedication to impactful problem-solving using data and machine learning, producing high-quality, maintainable code and achieving measurable business outcomes.

Work Experience

HumanWare Technologies Inc., *ML Engineer (Co-op)* | Montreal, Canada   *May 2024 - Dec 2024*

Skills/Tools: *PyTorch, LVLM (Fine-Tuning, Evaluation, and Benchmark), RAG, Prompting, LangChain, LangGraph, NLP (HunggingFace Transformers), CV (Object detection and Navigation), Distributed Training (Ray and vLLM Kernels), Linux, MLOps (AWS Bedrock)*

- Designed and implemented ETL pipelines to process and prepare 3TB of unstructured image-text multimodal data for fine-tuning a Large Vision-Language Model (LVLM).
- Developed an innovative Navigation System leveraging the fine-tuned LVLM and Agentic GenAI workflows to assist visually impaired users in navigating complex environments.
- Achieved a 32% improvement in navigation accuracy compared to the existing solution, alongside a $3\times$ reduction in system inference latency, driving a seamless and efficient user experience.
- Successfully deployed the devised method as a live product feature that resulted in a significant jump of 25% in retention rate.

Kumo.ai, *Machine Learning Developer* | Remote   *Mar 2023 - Aug 2023*

Skills/Tools: *PyTorch, PyTorch Geometric (PyG), Graph Neural Networks (GNNs) Matplotlib, Scikit-learn*

- Developed and integrated new live features and operators into Kumo's PyG toolkit such as Relational Graph Attention Networks, GNN Clustering for Graph-structured Data, and Model-agnostic GNN Explainability Framework.
- Designed and implemented unit tests to validate devised techniques' accuracy and efficiency.

IBM Systems Development Lab, *Senior Data Scientist* | Pune, India   *Apr 2022 - Mar 2023*

Skills/Tools: *TensorFlow, Time-Series Forecasting, SQL, AWS (DVC, EC2, Lambda, S3, CI/CD, Cloud Watch, SNS, API Gateway), Kubernetes (AWS EKS and KubeFlow), Apache Spark, MLFlow, Git, Problem-Solving, MLOps, Research Skills*

- Created a multivariate time series-based neural learning framework that utilizes generic machine learning algorithms (LSTM, GRU, and RNN) to impeccably predict run-out memory times of 300k+ storage systems in the resources pool, which resulted in the reduction of resource exhaustion time by 91%.
- Developed a novel approach to segment storage systems based on their respective predicted behaviors in terms of average growth in usable capacity over a specific forecasting window period, which in turn helped fill the missing data present in the time series.
- Designed and implemented pipelines for various stages of development, starting from extracting high-dimensional data in the form of multivariate time series (with underlying non-continuous dynamic) from a storage data lake to training ML time-series algorithms and predicting the storage systems' capacities.
- Created and optimized online inference approach to quickly serve the incoming requests for storage systems' capacity predictions at scale, leading to a 70% decrease in the framework's time to respond to user traffic requests.
- Introduced the entire framework as a fully operational machine learning-driven IBM Storage Insights product feature with low latency and high throughput, serving numerous customers' requests in real-time.

IBM Research Europe, *ML Research Software Engineer* | Rüschlikon, Switzerland   *Sept 2021 - Mar 2023*

Skills/Tools: *PyTorch, NLP, Anomaly Detection, PostgreSQL, AWS (DVC, EC2, Lambda, S3, SageMaker, CI/CD, Cloud Watch, SNS, API Gateway), Kubernetes (AWS EKS and KubeFlow), ML Explainability, HuggingFace, MLFlow, Apache Spark & Kafka, Model Compression & Quantization, Problem-Solving, MLOps, Research & Leadership Skills*

- Developed a novel unsupervised contrastive learning-based approach for the detection of similar bugs on 14TB of unstructured bug text corpus to facilitate faster log analysis.
- The approach extracts documents from the corpus of bug text and then implements Text Variational Autoencoders (T-VAEs)-based unsupervised anomaly detection technique to detect anomalies in the bugs corpus. Subsequently, learns a latent representation using the BERT language model to form a document embedding for each bug document, and then measures the degree of similarity between these learned document embeddings of all extracted documents using novel Word Centroids (WC) and Word Mover Distance (WMD) algorithms.
- Introduced quantization techniques to optimize this transformer-based similarity detection method, achieving a 38% smaller and $1.8\times$ faster model and attaining $4\times$ inference speedup and 14% improvement on both CPU and GPU.

- Upon obtaining the semantic-similarity scores, the approach clusters the similar bugs together based on semantic textual similarity indexing, which is governed by these computed scores.
- Included main measures to evaluate the approach (e.g. novel and repeated). Experiments conducted on the bugs corpus strongly recommended the 100% elimination of manual detection of similar bugs. This transition not only resulted in an impressive 85% reduction in response time to address client-reported bugs but also empowered the support team to focus solely on handling new and unique bug instances.
- Rolled out the devised method as a real-time IBM Storage Insights product feature functionality, accessible to a broad customer spectrum.

IBM Research India, *ML Research Software Engineer* | Bangalore, India   *July 2020 - Dec 2022*

Skills/Tools: PyTorch, PyTorch Geometric, Graph Neural Networks, Classification & Clustering for Graph Data, NetworkX, IBM Cloud Pak for Data (CP4D), Linux, Git, REST API, Problem-Solving, Research & Brainstorming Skills, Critical Thinking

- Developed a framework that converts property graphs of domain-specific ontologies into large-scale enterprise heterogeneous knowledge graphs, leveraging the multi-layered *RGCN* model for entity label prediction. This solution reduced entity duplication by 40%, improved data accuracy by 20%, and delivered $100K in annual savings on data management costs.
- Proposed and deployed an automated approach to evaluate GNN explanations, enhancing the reliability of predictions on graph-structured data. Work published at the ICML 2021 Explainable AI Workshop and integrated as a live feature in IBM AI Explainability 360 toolkit under the IBM Research Trusted AI initiative.
- Filed 2 patents: (i) Secure graph learning for classifying data ontologies; (ii) Faster deduplication of records to improve customer-specific assets in Data Fabric.

IBM Global Business Services, *Data Scientist* | Singapore   *Sept 2018 - July 2020*

Skills/Tools: PyTorch, TensorFlow, Scikit-learn, Pandas, Numpy, SQL, Azure (CI/CD Pipeline, Kubernetes Service with KubeFlow, Databricks, Serverless, Synapse Analytics), ML Explainability, Regression, ETL Processes, Git, Problem-Solving, Research & Brainstorming Skills, Critical Thinking, MLOps & Deployment

- Built a machine learning platform from scratch for monitoring and tracking the traversed path of 120k vessels.
- Wrangled 1.5TB of vessel trading data stored in the Synapse database to train a Machine Learning regressive prediction algorithm and implemented an offline batch inference method for vessel trajectory prediction (in terms of longitudes and latitudes), which reduced the overall vessel traversal logistics time by 27%.
- Led a cost-saving initiative by devising an automated ML approach to predict the Time Charter (TC) Rate to identify the best vessel at the lowest price based on the availability of the vessels, leading to an annual cost reduction exceeding $70,000 in logistics expenditures.
- Optimized the developed ML platform by implementing an explainability solution to identify influential features in the data for the prediction algorithm, resulting in improved trained ML model accuracy from 65% to 86% for both TC Rate and vessel trajectory prediction tasks.
- Integrated a data visualization framework, built using the *Dash-Plotly* library, into the platform for creating intuitive as well as user-interactive visualizations and launched the developed machine learning platform for production use, allowing 2k+ end clients to discover, on the fly, daily patterns in vessels movement and track the regular impact of vessel trades on liquidity.

Publications

- *V. BK, B. Ganesan, **A. Saxena**, D. Sharma, A. Agarwal, "Towards Automated Evaluation of Explanations in Graph Neural Networks," 2021 International Conference on Machine Learning (ICML) Workshop on Explainable AI (XAI).*
- ***A. Saxena** and A. Saxena, "Optimal Partition Search," 2019 IEEE International Conference on Electrical, Computer and Communication Technologies (ICECCT).*
- *A. Saxena, **A. Saxena**, J. Patel, "DeepCoder: An Approach to Write Programs," 2017 International Conference on Advanced Research and Innovation in Engineering (ICARIE).*

Education

- **B.Tech in Computer Science**, *Dr. A. P. J. Abdul Kalam Technical University | First Div. (Hons.)*   Aug 2014 - July 2018
- **MS in Computer Science and Machine Learning**, *Mila- Quebec AI Institute, University of Montreal*   Aug 2023 - Apr 2025
  Courses: Machine Learning, Data Science, Deep Learning, Natural Language Processing (NLP), Computer Vision

Achievements and Extracurricular Activities

- **2023 IBM Research Invention Achievement Award**.
- Awarded **Academic Excellence Award** for exceptional performance during Academic Year 2015-2016.
- **Service Excellence Award** by IBM Systems Development Lab for the year 2022.
- **University of Montreal - Academic Excellence Scholarship**.